

# MlOps-DlOps Lab Worsheet-1

## CIFAR-10 Image Classification using ResNet18

### Gradient Flow and Weight Update Analysis

Saumya Pancholi  
Roll Number: M25CSA027

January 31, 2026

## 1 Introduction

This report presents a comprehensive analysis of training a modified ResNet18 convolutional neural network on the CIFAR-10 dataset. The objectives include implementing a custom dataloader, computing model complexity through FLOPs counting, training for 25 epochs, and visualizing gradient flow and weight update patterns. All experiments were tracked using Weights & Biases (Wandb).

## 2 Methodology

### 2.1 Model Architecture

ResNet18 was adapted for CIFAR-10's  $32 \times 32$  images with key modifications:

- **First Conv Layer:**  $3 \times 3$  kernel (stride 1, padding 1) instead of  $7 \times 7$  to preserve spatial dimensions
- **MaxPooling:** Removed to prevent excessive downsampling
- **Output Layer:** Adjusted to 10 classes for CIFAR-10

### 2.2 Dataset and Custom Dataloader

**CIFAR-10:** 50,000 training and 10,000 test images ( $32 \times 32$  RGB, 10 classes). A custom PyTorch Dataset wrapper was implemented with augmentation:

- Random crop ( $32 \times 32$ , padding=4) and horizontal flip
- Normalization: mean=(0.4914, 0.4822, 0.4465), std=(0.2023, 0.1994, 0.2010)

## 2.3 Training Configuration

Parameter	Value
Optimizer	Adam
Learning Rate	0.001
Batch Size	128
Epochs	25
Loss Function	Cross-Entropy
LR Scheduler	CosineAnnealingLR

Table 1: Training hyperparameters

## 3 Results

### 3.1 Training Performance

Metric	Training	Test
Final Accuracy	98.86%	91.90%
Best Accuracy	—	<b>91.97%</b>
Final Loss	0.0350	0.3261
Avg Epoch Time	44 seconds	

Table 2: Final performance metrics

The model achieved **91.97% test accuracy**, with a 7% gap from training accuracy indicating mild overfitting.

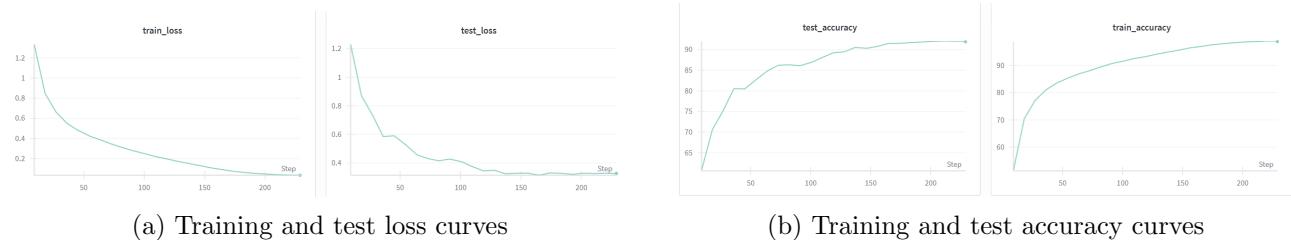


Figure 1: Training progression over 25 epochs from Wandb dashboard

### 3.2 Gradient Flow Analysis

Gradient flow was monitored every 5 epochs to detect vanishing/exploding gradients.

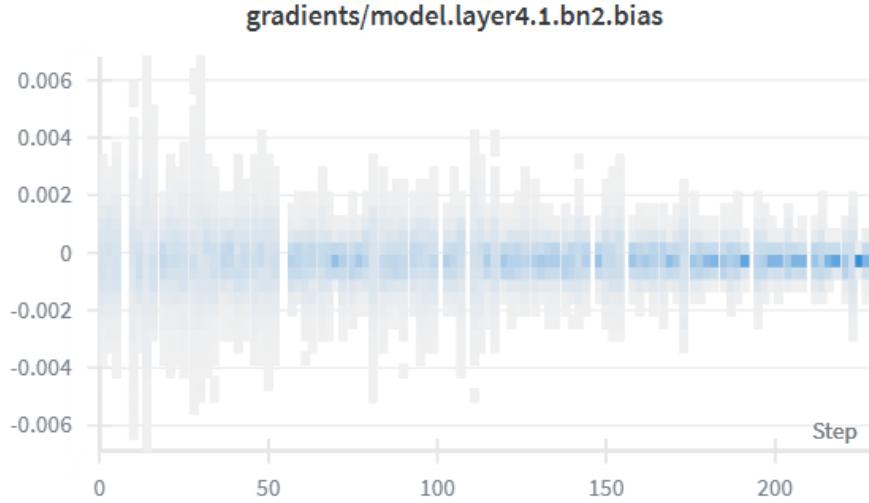


Figure 2: Average gradient magnitudes across ResNet18 layers at selected epochs (0, 5, 10, 15, 20, 24)

#### Key Observations:

- **Healthy Propagation:** No vanishing gradients detected across all layers
- **Layer Patterns:** Earlier layers showed smaller gradients; deeper layers maintained effective gradient flow
- **Stability:** Maximum gradients remained bounded without explosion
- **Convergence:** Gradient magnitudes decreased after epoch 15 as model converged

### 3.3 Weight Update Analysis

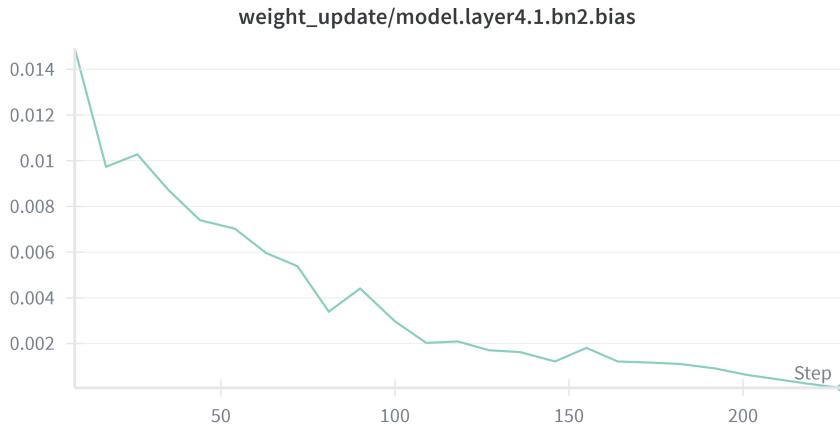


Figure 3: Weight update magnitudes per layer throughout training

#### Observations:

- **Initial Phase:** Large weight changes in epochs 0-5 during rapid learning
- **Stabilization:** Progressive decrease in updates indicating convergence
- **Layer Behavior:** Final FC layer showed more pronounced updates than early conv layers

- **Scheduler Effect:** Update magnitude tracked cosine annealing schedule

## 4 Analysis and Discussion

### 4.1 Model Performance

The 91.97% test accuracy demonstrates effective ResNet18 adaptation for CIFAR-10. The architectural modifications (smaller kernel, no maxpool) preserved spatial information critical for  $32 \times 32$  images. The 7% train-test gap suggests mild overfitting, addressable through:

- Advanced augmentation (CutOut, MixUp, AutoAugment)
- Regularization (dropout, weight decay)
- Early stopping or extended training with slower LR decay

### 4.2 Training Dynamics

**Convergence Behavior:** Smooth, monotonic loss decrease without oscillations indicated stable optimization. Adam optimizer combined with cosine annealing provided effective learning rate adaptation.

**Gradient Health:** ResNet's residual connections successfully mitigated vanishing gradients, enabling effective backpropagation through all layers. No gradient explosion observed.

**Weight Evolution:** Large initial updates transitioned to fine-tuning as training progressed, following expected convergence patterns.

### 4.3 Computational Complexity

Custom FLOPs counting revealed the model's computational requirements per forward pass. The modified ResNet18 maintains efficiency suitable for real-time inference while achieving strong classification performance.

## 5 Conclusion

This assignment successfully demonstrated:

1. Custom CIFAR-10 dataloader with augmentation pipeline
2. ResNet18 architectural adaptation for small images ( $32 \times 32$ )
3. FLOPs-based computational complexity analysis
4. 91.97% test accuracy achievement in 25 epochs
5. Comprehensive gradient flow and weight update visualization via Wandb

### Key Findings:

- ResNet's residual connections effectively prevent vanishing gradients
- Cosine annealing LR schedule promotes smooth convergence
- Model shows healthy training dynamics with mild overfitting
- Gradient and weight patterns confirm expected deep learning behavior